Problem-Oriented Analysis of Self-Organizing Neural Networks

Vins Sharma
Electrical and Computer Engineering
Carnegie Mellon University
vmsharma@andrew.cmu.edu

Anand Raju
Electrical and Computer Engineering
Carnegie Mellon University
amraju@andrew.cmu.edu

***Abstract*—Liquid state machines (LSMs) are machine learning models that create more biologically plausible neural networks by self-organizing neurons into a 'liquid'. As such, we aim to analyze a variety of 'liquids' generated off unique problems and datasets in order to determine correlations and patterns that exist in these graphs. We accomplish this by: (1) Implementing a framework to generate LSMs on many problems, and (2) Analyze generated LSMs for overlapping patterns and similarities.**

***Index Terms*—liquid state machines, temporal neural networks, random graph models, reservoir computing**

I. Introduction

A. Motivation

The portion of the brain primarily responsible for human thought and cognitive function is thought to be the neocortex, which covers the outside shell. The unfolded neocortex is the size of a dinner napkin, and is made up of tiny perpendicular 'micro-fibers' called cortical columns [1].

Cortical columns are thought to be the microcircuit that implements all cognitive thought. Though it is well documented that different regions of the brain are responsible for different functions, it has been shown [2] that the makeup of the cortical column is nearly identical across the entire neocortex, excluding the inputs. This indicates that a circuit exists that can be plugged into different problems and solve all of them without significant modification or significant power consumption.

One approach to designing such a 'perfect' circuit is through induction. In section II, we'll discuss this approach in more detail, and why we are not choosing it. However, another approach is from a top down - We hope to generate networks that can solve a variety of problems, look at these networks, and attempt to understand what similarities exist between them or how different parameters and problem-specific features affect their design.

B. General Approach

We start by generating a framework to quickly and easily 'plug in' different datasets and generate networks. These networks can organize themselves based on features of the data, and we can analyze how they organize themselves over time.

Once we have a framework that can handle this, we aim to introduce a series of unique problems and analyze how individual neurons in a network may connect. Ideally, we can draw significant conclusions about the influence of various parameters, and maybe even find similarities. These similarities and conclusions can serve as the first steps in constructing a cortical column-like neural circuit.

II. Previous Work

In this section we examine inspirations for our approach, and potential choices.

A. Smith (2023)

As discussed in I, the goal of this project is to get closer to a cortical column-like neural circuit. Professor Smith's rendition of his macrocolumn in [4] is an example of the ground-up approach to solving this problem. He proposes a circuit specialized towards one task - Having a mouse navigate a maze. Future work in this project involves applying a slightly modified circuit to another problem.

One major issue with this avenue of thought is the exponentially increasing amount of effort. In future work, when introducing a new problem, Professor Smith now has to modify the circuit to solve the new problem while still preserving the old functionality. As such, we aim to avoid this approach, and instead take a top-down approach to generalize further.

B. Nair, Shen, Smith (2021)

This paper [3] describes the temporal neuron and a strategy for training it in an unsupervised way. The temporal neuron simply encodes data inputs on spike timing lines (in a process known as rate coding). This allows it to consume significantly less power (as the spike is only on instantaneously), and also allows for a non-statistical simple unsupervised approach.

For our approach, we implement these temporal neurons, but use them without the paper's described 'columnar' structure (for the most part). As such, our usage of temporal neurons is novelly applied in a more arbitrary network structure.

C. Purdy (2016)

Purdy's paper [5] discusses the best-practices for encoding data in temporal systems, such as we will throughout this project. He focuses on the generation of proper sparse distributed representation (SDR) matrices as good-for-encoding inputs. We apply this knowledge to adapt datasets to temporal input.

D. Maass (2011)

Maass defines the liquid state machine (LSM) [6] model for reservoir computing. The LSM has a reservoir with a series of spiking neurons that form some arbitrary excitation pattern; A separate set of (trained) neurons then read out this pattern and correlate it to an output. The internals of the reservoir can be structured in any way, but are conceptually similar to self-organizing maps [7].

Kohonen posits, through his work, that arbitrary higher dimensional features of a problem are encoded into the graph's organization for a self-organizing map. As such, we analyze the organization of the reservoirs in order to draw possible conclusions about the features of the problem.

E. Hazan, Manevitz (2012)

Hazan and Manevitz [8] propose in their paper that the LSM previously defined by Maass is not a good model due to its network's robustness, and that a more robust model would better represent the brain. They draw a new conclusion towards small-world assumptions in graph topology being key to solving this problem. We re-test this hypothesis in our work.

III. Approach

A. Datasets

Throughout our project, we apply a few datasets and analyze their generated graphs. The datasets we use are, relatively speaking, orthogonal to each other - They solve fundamentally different problems. As such, any similarities may be in response to parameter changes, the network structure, or 'problem-solving' itself.

Our first benchmark is the MNIST dataset [9]. This dataset requires the tools of image recognition and classification, and is popular in machine learning projects.

From there, we look at the Numenta Anomaly Benchmark dataset [10]. Anomaly detection is a problem that temporal neural networks generally specialize in, so we analyze these as well.

Finally, we generate and analyze a network for the NFL Big Data Bowl dataset [11]. This network would perform prediction for the next 'play' (e.g. pass or rush) based on the previous play's characteristics.

B. Network Configuration

Before discussing our network, we introduce our system. We define a temporal liquid state machine (TLSM) as an input vector, encoded in time-based spikes, followed by a reservoir of temporal neurons, and finally a singular output column (with winner-takes-all lateral inhibition). The input vector uses encoding strategies discussed in [5], and the output column is identical to the one described in [3].

The reservoir is the focus of network analysis, containing a series of neurons. This graph of neurons will change over time with training and with STDP rules, adding and removing connections arbitrarily in response to the inputs. The neuron activation function is also a hyperparameter of the network.

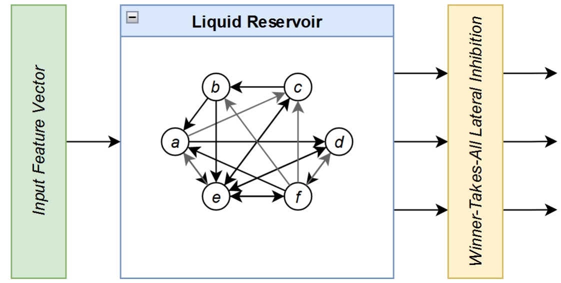

Fig. III.1. Diagram of a temporal liquid state machine (TLSM).

Furthermore, a 'seed' configuration for the network may be specified, indicating its initial connectivity. This seed configuration is generated through a number of random graph algorithms, and is a hyperparameter of the network.

C. Assumptions

While we vary a variety of parameters (such as neuron activation function), we assume that the underlying temporal neurons function as a strong analogue for biological processes. This assumption allows us to generalize our results to neural basis of cognition rather than design of artificial neural networks.

D. Analysis

The focus of analysis is the network generated within the liquid reservoir. This liquid has nodes of neurons and edges of weighted synapses.

We test the robustness claim of [8] by analyzing the connectivity of these networks over their growth. We also analyze diameter in order to understand the 'distance' between neurons in the network (and the amount of time it takes for a spike to propagate).

E. Research Questions

Through these network features, and more, we aim to understand the following:

1) Why did the network configure itself the way it did?
2) How do networks of different problems compare to each other?

IV. Experimental Setup and Results

Our first benchmark and first goal is the MNIST dataset [9] and the creation of a framework that can automate this procedure. In this milestone we have completed the design for the Step No-Leak (SNL) neuron, as well as the input vector and liquid reservoir. Other neuron types are in progress, and the temporal mini-column enabling supervised training has yet to be implemented.

We also have a number of seed networks implemented and ready for initializing and training weights - These include Erdos-Renyi, Watts-Strogatz, and Barabasi- Albert random graphs. We also have a fully-connected structure for comparison.

V. Conclusions

Our project presents a framework for generating and analyzing self-organizing neural networks in order to derive insight into the neural basis for cognition through the orientation of the problem. Through analyzing the networks, we hope to uncover similarities and understand how these structures may organize and train to particular problems.

A. Plans for Next Milestone

By the end of October, we plan to complete our network generation framework, alongside a number of other seed networks and neuron types. We also plan to generate and analyze the generated MNIST network [9].

B. Contribution

For this milestone, Vins focused on the input vector and SNL neuron implementation. In contrast, Anand focused on the initial graph organization and seed matrix generation. Both teammates contributed to the overall design of the TLSM generator framework, and both teammates will contribute to design and analysis of the MNIST network [9].

In the future, Anand will be responsible for the design and analysis of an NFL data network [11], and Vins will be responsible for the NAB network [10].

References

[1] V. B. Mountcastle. “The Columnar Organization of the Neocortex.” *Brain: A Journal of Neurology* (1997).

[2] J. Hawkins, R. Dawkins. “A Thousand Brains: A New Theory of Intelligence.” *Basic Books* (2021).

[3] H. Nair, J. P. Shen, J. E. Smith. “A Microarchitecture Implementation Framework for Online Learning with Temporal Neural Networks.” *IEEE Computer Society Annual Symposium on VLSI (ISVLSI)* (2021).

[4] J. E. Smith. “A Macrocolumn Implemented with Spiking Neurons.” *arXiv preprint arXiv:2207.05081* (2023).

[5] S. Purdy. “Encoding Data for HTM Systems.” *arXiv preprint arXiv:1602.05925* (2016).

[6] W. Maass. “Liquid State Machines: Motivation, Theory, and Applications.” *Computability in Context: Computation and Logic in the Real World* (2011).

[7] T. Kohonen. “Self-Organized Formation of Topologically Correct Feature Maps.” *Biological Cybernetics* (1982).

[8] H. Hazan, L. M. Manevitz. “Topological Constraints and Robustness in Liquid State Machines.” *Expert Systems with Applications* (2012).

[9] AstroDave, W. Cukierski. “Digit Recognizer.” *Kaggle* (2012). www.kaggle.com/competitions/digit-recognizer

[10] Numenta. “Numenta Anomaly Benchmark (NAB).” *Kaggle* (2016). www.kaggle.com/boltzmannbrain/nab

[11] The National Football League. “NFL Big Data Bowl.” *Kaggle* (2022). www.kaggle.com/competitions/nfl-big-data-bowl-2022